



Between-reading variance predicts which ambiguous items benefit from clarification

Qwen3-8B (n = 30 Set A items)



Llama-3-70B (n = 44 Set A items)



Points plot the full 32-sample estimates. The reported correlation is the split-half one: the predictor is estimated from 16 draws per prompt and the outcome from the disjoint 16, so the two share no sampling noise. Both half-assignments are shown.